

Big Data - Labo

Tuur Vanhoutte, Jakob Soens

May 4, 2021

Contents

1	Intro	1
2	NAT-ing	1
2.1	NAT	1
2.1.1	The problem	1
2.1.2	The solution	2
2.2	SSH Tunnel	2
3	Container technology	2
3.1	Docker	2
3.2	Microservices	3
3.2.1	Monolithic vs Microservices	3
3.3	Virtualization vs Containerization	4
3.3.1	Virtualization	4
3.3.2	Containerization	4
3.4	Shared kernel	4
3.4.1	What is a kernel?	4
3.4.2	Namespaces	6
3.4.3	Containers	6
3.5	Images	7
3.5.1	Image layer	7
3.6	Docker is lightweight	7
3.7	Using Docker	8
3.7.1	View layers	8
3.7.2	Make Dockerfile more efficient	8
3.7.3	Connecting to a database in a different container	9
4	Sharding	9
4.1	Create index	9
4.2	Health	10
4.2.1	Shard health	10
4.2.2	Index health	10
4.2.3	Cluster health	10
4.3	Shard allocation	10
4.3.1	Unassigned	10
4.3.2	Initializing	11
4.3.3	Started	11
4.3.4	Relocating	12
4.4	Change number of replicas	12
4.4.1	Health when one fails	13
4.5	Caveat: single node cluster	13
5	Linux batching + Dask	13
5.1	Python & data engineering/science	13
5.2	Spark vs Dask	13
5.2.1	Spark	13
5.2.2	Dask	14
6	TICK Stack	14
6.1	Telegraf	14

6.2	InfluxDB	14
6.2.1	Key concepts	15
6.3	Chronograf	15
6.4	Kapacitor	16
6.5	Deployment models	16
6.6	Architecture of the TICK stack	17
6.6.1	Write Ahead Log (WAL)	17
6.6.2	Time Structured Merge Tree (TSM)	17
6.6.3	Time series index (TSI)	18
6.6.4	Sharding	18
6.7	Pitfalls, tips & tricks	19
7	Mogelijke examenvragen	19

1 Intro

Topics:

- Linux basics + containers
- Elastic search (text search, document store)
- Linux Batch Processing & Dask
- InfluxDB (timeseries)
- Cloud services (Kafka, Kinesis, Lambda, ML services, ...)

2 NAT-ing

2.1 NAT

= Network Address Translation

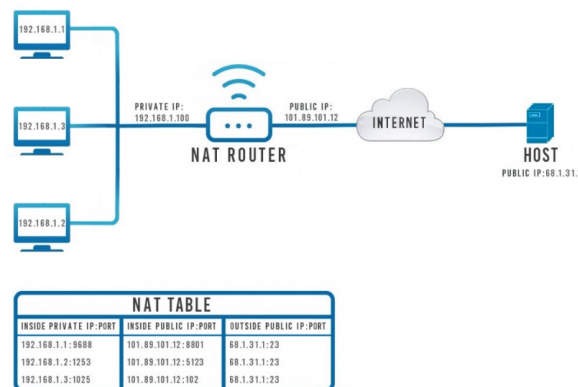


Figure 1: NAT diagram

2.1.1 The problem

- We only have one (public/private) IP-address
 - Howest: 172.23.82.60
- Connecting to a server over a network:
 - Using a protocol (HTTP) which uses TCP
 - Our server has an IP address: 172.23.82.60
 - Our server is listening at port 5000
 - \Rightarrow `http://172.23.82.60:5000`
- Problem: We want to have multiple IP addresses
 - Student 1 wants to reach `http://192.168.20.21:5000`
 - Student 2 wants to reach `http://192.168.20.22:5000`
 - Student x wants to reach `http://192.168.20.xx:5000`

2.1.2 The solution

Translation is needed!

- 172.23.82.60:5000 should point to 192.168.20.21:5000
- 172.23.82.60:5001 should point to 192.168.20.22:5000
- 172.23.82.60:5xxx should point to 192.168.20.xx:5000

We can use any port, on both sides:

- 172.23.82.60:8000 can point to 192.168.20.21:5000
- 172.23.82.60:8000 can point to 192.168.20.21:3000

2.2 SSH Tunnel

= SSH Port Forwarding

Resource	Internal IP	Username	Password	External port	Internal port
Vyos Router	192.168.50.1	vyos	P@ssw0rd	7000	22
Storage	192.168.50.2	student	P@ssword	n.v.t.	22
SSH	192.168.50.3	student	P@ssword	7040	22
RDP	192.168.50.4	Administrator	P@ssword	7020	3389
vCenter vSphere	192.168.50.10	administrator@vsphere.local	P@ssword	7060	443
vCenter appliance	192.168.50.10	root	P@ssword	n.v.t.	5480
ESXi-00	192.168.50.11	root	P@ssword	n.v.t.	22
ESXi-01	192.168.50.12	root	P@ssword	n.v.t.	22

Figure 2: Example



Figure 3: Example: a tunnel is opened and we log into user@instance

3 Container technology

3.1 Docker

- Docker = ecosystem for creating and running containers
- Docker wants to make it possible to install and run software on any system

- Other reasons: Microservices/DevOps/Resource usage
- Docker != Container
 - Docker CLI
 - Docker Engine
 - Docker Image
 - Docker Container
 - Docker Hub
 - Docker Compose
 - Docker Swarm
 - ...

3.2 Microservices

- = A software development technique
- Structure an application as a collection of loosely coupled services
- Lightweight
- Microservices-based architectures enable continuous delivery and deployment
- <https://en.wikipedia.org/wiki/Microservices>

3.2.1 Monolithic vs Microservices

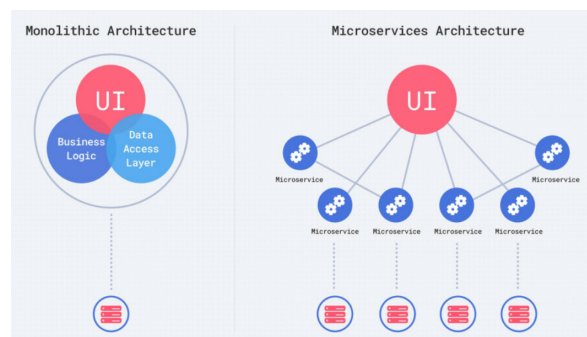


Figure 4: Monolithic architecture vs Microservices architecture

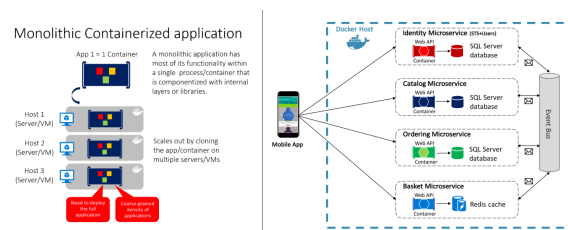


Figure 5: Monolithic Containerized application

Microservices does **not** necessarily mean containerization!

3.3 Virtualization vs Containerization

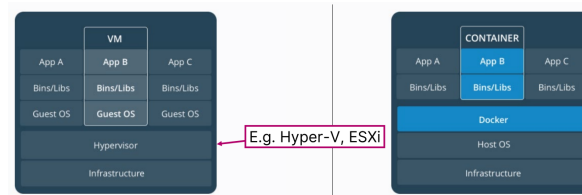


Figure 6: Virtualization vs Containerization

3.3.1 Virtualization

- = An abstraction of physical hardware turning one server into many servers
- Multiple VMs can run on the same machine
- Each VM includes a full copy of an Operating System (OS), one or more apps
- Takes a lot of space
- Can be slow to boot

3.3.2 Containerization

- = An abstraction at the app layer that packages code and dependencies together
- Multiple containers can run on the same machine, they share the OS kernel with each other, each running as isolated processes in user space.
- Takes up less space than VMs
- Boot up almost instantly

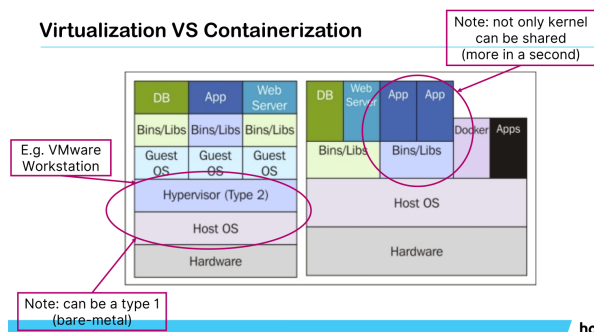


Figure 7: Schematic

3.4 Shared kernel

3.4.1 What is a kernel?

- Piece of software that offers basic functionality to the OS

- System calls: open, read, write, close, wait, exit, . . .
- A typical kernel has a few hundred system calls

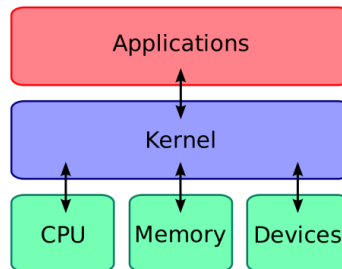


Figure 8: The kernel is the layer that communicates between hardware and applications

- Docker shares the host OS kernel
 - Host OS: Windows / MacOS / Linux
 - Shared Linux Kernel

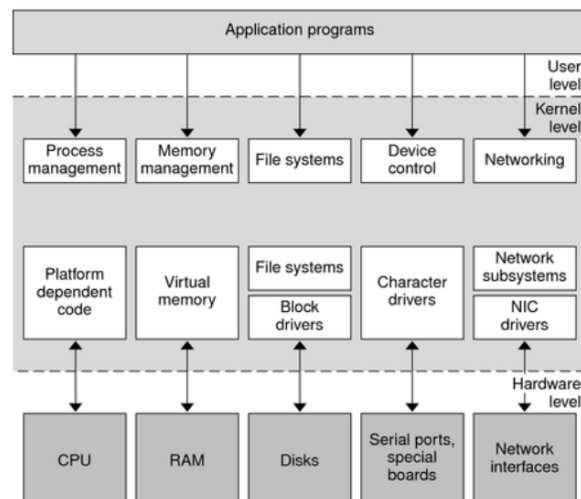


Figure 9: Kernel in detail

- The Ubuntu container requires the Linux kernel
- The Linux kernel runs in a Virtual Machine

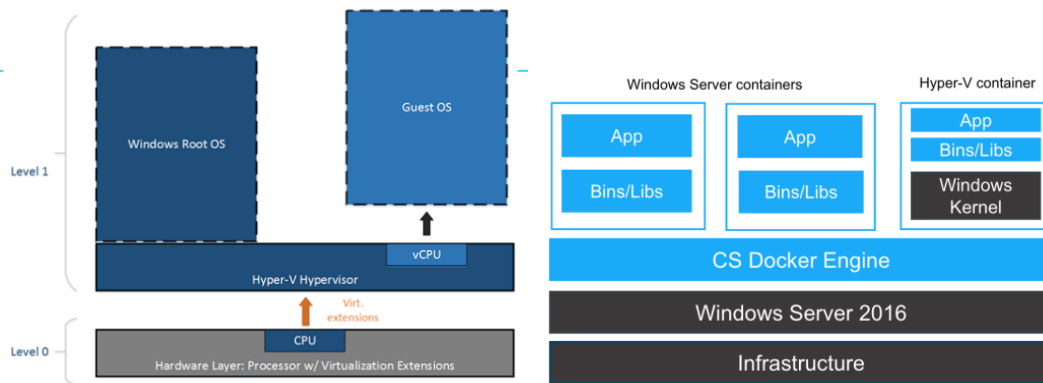


Figure 10

3.4.2 Namespaces

Docker uses two important Linux kernel features to share the kernel:

- **Namespaces** are a feature of the Linux kernel that partitions kernel resources
- **cgroups** (control groups) is a Linux kernel feature that limits, accounts for, and isolates resource usage of a collection of processes

Simpler:

- Namespaces = isolating resources per process (or group of processes)
- cgroups = a type of namespace that limits resource usage per process (or group of processes)

7 types of namespaces:

- mount, UTS, IPC, network, PID, cgroup, user
- For the process (or group of processes) it looks like there is a completely isolated set of resources

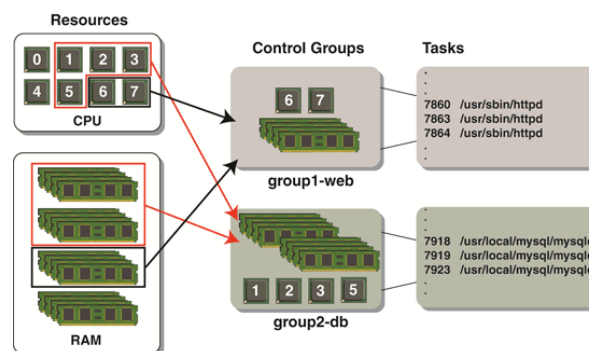


Figure 11: Control groups (cgroups)

3.4.3 Containers

What is a container?

- One or more running processes (if not running anymore \Rightarrow container dead)

- Resources are specifically assigned to it
- The real building blocks: Linux kernel features
 - Namespaces
 - cgroups

3.5 Images

What is an image?

- Filesystem snapshot
- Startup command
- Layered structure (!)

Instance of image = container

3.5.1 Image layer

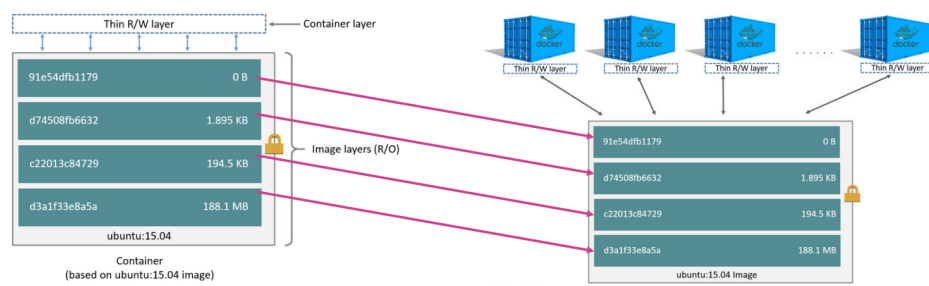


Figure 12: Image layers

- RUN, COPY, ADD
 - = new read-only layer
- Top layer = container layer
 - Writeable
- Delete container = delete container layer
 - Image will still exist
 - Persistent volumes

3.6 Docker is lightweight

- Shared kernel
- Container has no OS
- Less disk space \Rightarrow sharing layers
- Small community images

- ex: Alpine Linux (small, simple, secure)
- Current Docker version is using runC (previously LXC = Linux Containers)
 - runC = tooling (written in Go) that makes it possible to create and run containers
 - runC = CLI to ‘easily’ access kernel features such as cgroups and namespacing
 - runC = successor of libcontainer (developed by Docker)
 - Open-sourced ⇒ better community
 - runC implements ‘Open Container Initiative Runtime Specification’
 - <https://github.com/opencontainers/runtime-spec>

Docker is ‘nothing more’ than an ecosystem about creating & running containers

3.7 Using Docker

(see slides 40-55 in [02_big_data_01_containers.pdf](#) for basic commands)

3.7.1 View layers

With the command ‘docker history <image | container id>’ you’ll get an overview of the layers of an image.

- Every RUN, COPY, ADD adds a new read-only layer
- Make Dockerfile more efficient ⇒ create less layers

3.7.2 Make Dockerfile more efficient

Our Dockerfile, before optimisation:

```

1 FROM python:3.9.1-alpine3.13
2 WORKDIR '/app'
3 RUN apk add --no-cache linux-headers g++
4 RUN pip install Flask # we can replace these two lines by:
5 RUN pip install uwsgi # RUN pip install -r requirements.txt
6 COPY ./ ./
7 RUN addgroup -S uwsgi && adduser -S uwsgi -G uwsgi
8 USER uwsgi
9 CMD ["uwsgi", "--ini", "app.ini"]

```

After optimisation:

```

1 FROM python:3.9.1-alpine3.13
2 WORKDIR '/app'
3 RUN apk add --no-cache linux-headers g++
4 # the addgroup and adduser commands can be higher up
5 RUN addgroup -S uwsgi && adduser -S uwsgi -G uwsgi
6 # first, we copy the requirements.txt file
7 COPY ./requirements.txt ./
8 # then we install ALL packages
9 RUN pip install -r requirements.txt
10 # then we copy the remaining files

```

```

11 COPY ./ ./
12 USER uwsgi
13 CMD ["uwsgi", "--ini", "app.ini"]

```

3.7.3 Connecting to a database in a different container

Use 'ip a' to find the correct ip to use in this command:

```

1 docker run -p 8080:8080
2   -e POSTGRES_PASSWORD=student_password
3   -e POSTGRES_USER=student_user
4   -e POSTGRES_DATABASE=labo
5   -e POSTGRES_PORT=5432
6   -e POSTGRES_HOST=ip-van-je-vm # change this ip
7   -e PORT=8080
8   jouw-naam/api                # change this

```

4 Sharding

- Index = collection of documents
- Document = data in JSON format
- Shard = A piece of an index. Index is "sharded" in blocks, a block = shard
- Primary shard = Document is primarily indexed (written) to a primary shard
- Replica shard = an asynchronous copy of the primary shard

4.1 Create index

```

1 {
2   "settings": {
3     "number_of_shards": 2,
4     "number_of_replicas": 1
5   }
6 }

```

How many shards in total: **4**

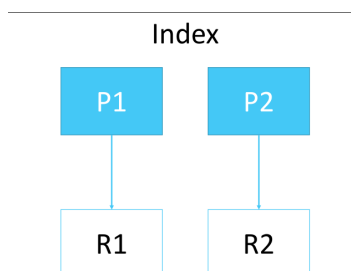


Figure 13: 4 shards: 2 primary shards with 1 replica each

4.2 Health

Health exists at shard, index and cluster level!

4.2.1 Shard health

- Green = all shards are allocated
- Yellow = all primaries are allocated but at least one replica is not
- Red = at least one primary shard is not allocated in the cluster

4.2.2 Index health

= status of the worst shard in that index

4.2.3 Cluster health

= status of the worst index in the cluster

4.3 Shard allocation

Shards states:

- Unassigned = master did not assign the shard (yet)
 - Or master is not able to assign the shard
- Initializing = master did assign the shard, creating...
- Started = shard is fully operational
- Relocating = shard is moving
 - Imbalance, new nodes, removed nodes, ...

4.3.1 Unassigned

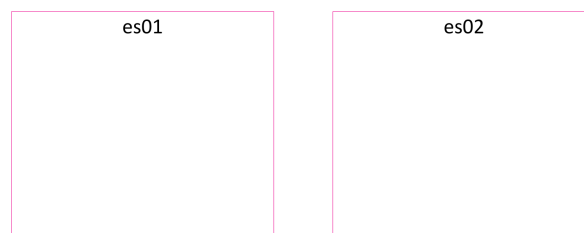


Figure 14: No shards assigned yet

4.3.2 Initializing

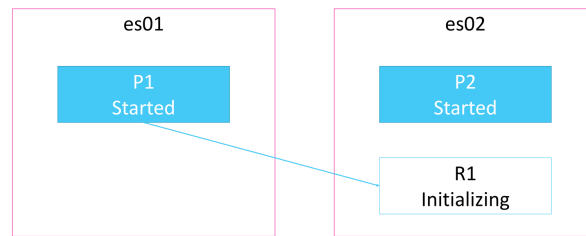


Figure 15: Creating shards

4.3.3 Started




Figure 16: The primary shards have been started, replica 1 is initializing. **Cluster status = yellow**

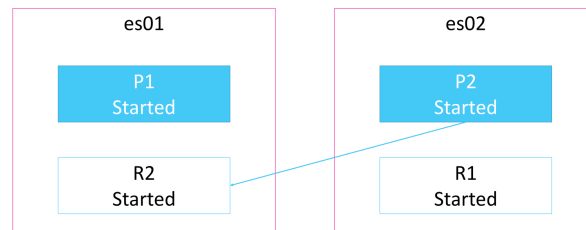


Figure 17: **Cluster status = green**

What if one of the node fails?

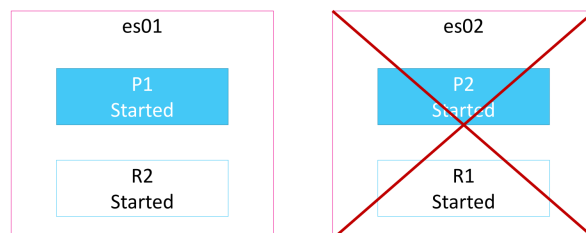


Figure 18: Situation when one node fails

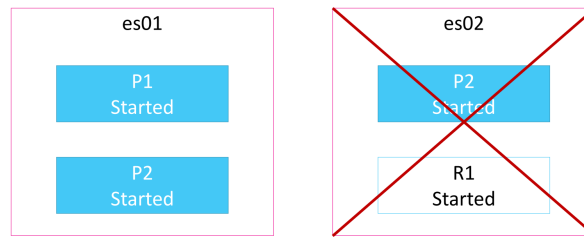


Figure 19: After some time, R2 will become a primary shard. **Cluster status = yellow**

4.3.4 Relocating

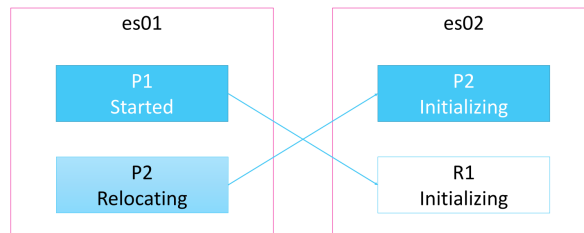


Figure 20: After es02 is restored, P2 gets relocated to its previous node

4.4 Change number of replicas

```

1 {
2   "index": {
3     "number_of_replicas": 0
4   }
5 }
```

How many shards in total? **2**

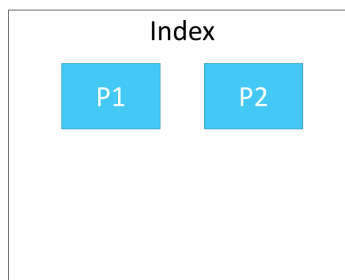


Figure 21: 2 shards total: 2 primary shards, 0 replicas each

4.4.1 Health when one fails




Figure 22: **Cluster status = Red**

4.5 Caveat: single node cluster

- Bootstrap checks = important settings are checked
- `discovery.type=single-node`
- If a node is already part of a cluster:
 - Unique node ID
 - Unique cluster ID
 - Not easy to create a new cluster

5 Linux batching + Dask

5.1 Python & data engineering/science

- Many tools, libraries (numpy, pandas)
- Not very scalable \Rightarrow parallelisation
- Threads/processes possible, but complex and not ideal
- What if it doesn't fit in memory?
 - To disk?
 - Possible, but complex! Some operations require everything in memory

5.2 Spark vs Dask

5.2.1 Spark

- Complex, learning curve!
- Complete 'engine', clustering
- Streaming engine
- Written in Java: uses the Java Virtual Machine (JVM) \Rightarrow not very accessible
- Standalone

5.2.2 Dask

- Simpler (especially if you know Python)
- Lightweight, even useful when only 1 node
- More flexible, but less performance
- Integration with other libraries
- 'The Python version of Spark'

6 TICK Stack

The TICK stack is an acronym for a platform of open source tools built to make collection, storage, graphing, and alerting on time series data incredibly easy.

The tools:

- Telegraf
- InfluxDB
- Chronograf
- Kapacitor

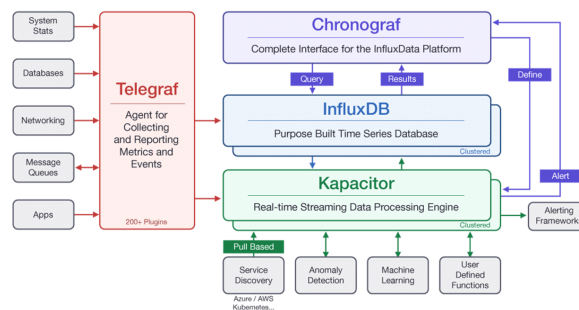


Figure 23: The components of the TICK Stack

6.1 Telegraf

- = Agent for collecting and reporting metrics and events
- Has inputs and outputs

6.2 InfluxDB

= Purpose built time series database

- Open source
- Simple HTTP API (POST, GET) with client libraries
- Somewhat similar to classic SQL, there are two versions:
 - V1: SQL & Flux: `SELECT * FROM measurement WHERE tag=value`

- V2: Flux, less like SQL, better for time series data:

```

1  # Flux
2  from(bucket: "bucket")
3  |> range(start: v.timeRangeStart, stop: v.timeRangeStop)
4  |> filter(fn: (r) => r["_measurement"] == "test")

```

6.2.1 Key concepts

- Line protocol = a text-based format that provides the measurement, tag set, field set, and timestamp of a data point:

```

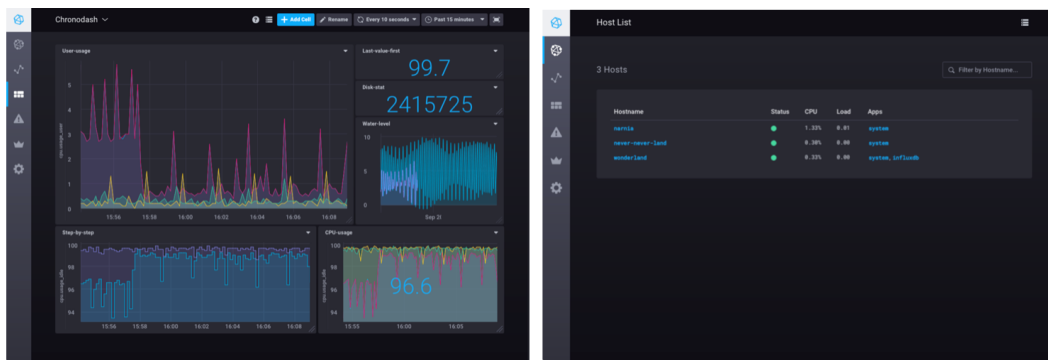
1  weather,location=us-midwest temperature=79,humidity=49 1591711854359
2  weather,location=us-midwest temperature=82,humidity=50 1591711787540
3
4  |           |           |           |
5  |           |           |           |
6  +-----+-----+-----+-----+-----+
7  |measurement|,tag_set| |field_set|           |timestamp|
8  +-----+-----+-----+-----+-----+

```

- Measurement = data that belongs together
- Timestamp = UNIX format
- Tags / Fields = key:value
- Tag = metadata
 - Tags are indexed
 - 'Fields' where you want to query on
 - Only strings!
- Field = data
 - Fields are not indexed
 - Floats, integers, strings, and booleans
- Tag set = set of tags
- Field set = set of fields

6.3 Chronograf

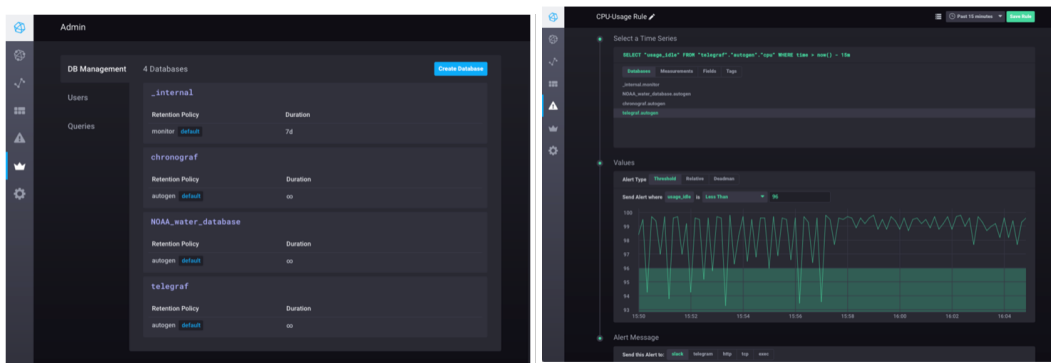
= A visualization tool



Data Visualization

Infrastructure Monitoring

Figure 24: Data visualization and Infrastructure Monitoring



Database Management

Alert Management

Figure 25: Database management and alert management

6.4 Kapacitor

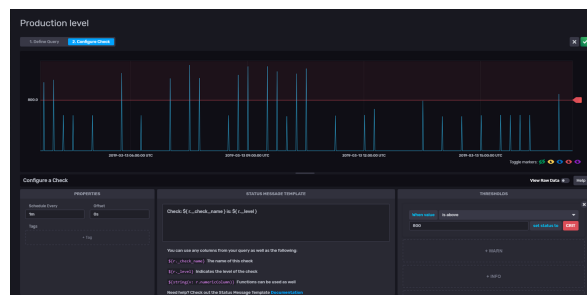


Figure 26

6.5 Deployment models

V2

- Open source version (OSS)

- No clustering
- No out-of-the-box replication
- Enterprise version: expensive, contact sales
- Cloud version: cheaper, usage based
- Chronograf and InfluxDB: one component
- Multi-tenant focus

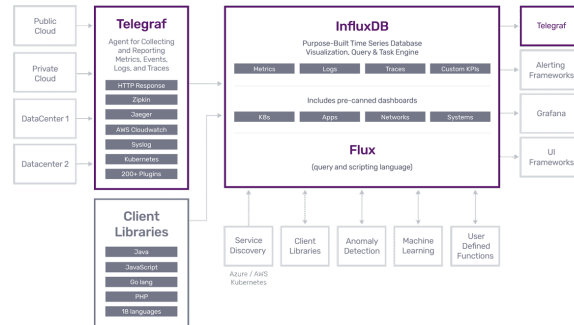



Figure 27: InfluxDB 2.0: a better graphic

6.6 Architecture of the TICK stack

6.6.1 Write Ahead Log (WAL)



- Disk optimized format (fast writes ↔ slow queries)
 - Not optimized for fast queries ⇒ memory!
- In case of a crash: replay WAL (durability ↑)
- What if we have more data than memory?
 - Out-of-memory errors (OOM)
- InfluxDB v1 vs v2 & OSS vs cloud:
 - V1 & V2 OSS: flat, simple file
 - V2 cloud: Kafka

6.6.2 Time Structured Merge Tree (TSM)

- A data structure optimized for storage and fast time-series queries

- Compressed data in columnar format
- Easy memory-mapping
- Similar to Log Structured Merge Tree (LSM)
- Field values are grouped by series key, ordered by time
- Series key = measurement, tag set and (a single) field key

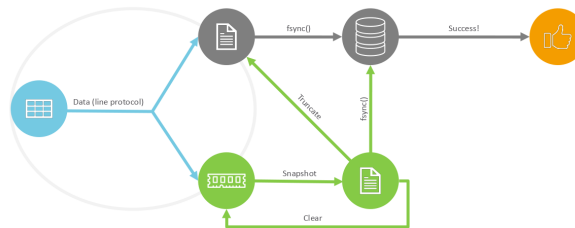


Figure 28: TSM + WAL

- Fast(er) queries: only read required series
- Compression: saved data is smaller than original (more data per node)
- Columnar: easy for memory-mapping
 - Data is cached for a limited time (solves OOM)
- What if we have many series keys (high cardinality)
 - Finding the right data will be slow!

6.6.3 Time series index (TSI)

- A data structure optimized for storage and fast query of series keys
 - TSM stores the data grouped by the series key
 - TSI stores the series keys grouped by measurement, tag and field key
- TSI answers two questions:
 - What measurements, tags and fields exists?
 - Given a measurement, tag, field, what series key exists?
- TSI stores the index in memory and on disk
 - Memory = page cache (least recently used memory)
 - Disk: writes to a WAL, compaction in the background

6.6.4 Sharding

V1:

- Directory with WAL, TSM and TSI files
- Retention policy (on database level)
- Each shard has a start and endtime

- Scalability ... but only for InfluxDB Cloud / InfluxDB Enterprise

V2:

- Sharding in V1 has much overhead: WAL, TSM and TSI / shard
 - Too much redundant data, especially for the TSI
 - Too many writes
- Not everyone needs a retention policy
- Sharding is now implemented as a block, like in most other database systems (in OSS only 1 shard)

6.7 Pitfalls, tips & tricks

- Tips for optimal (write) performance:
 - Order your timestamps
 - Order your tags alphabetically
 - Use the right precision: seconds, milliseconds, microseconds or nanoseconds
 - Write in bulks (less fsync's)
- Duplicates: measurement, tag set & timestamp
- Tags vs. Fields
- V2 is a great product, but:
 - Documentation is far from complete
 - Bugs in client libraries, e.g. precision is neglected
 - Quick release cycle / bug fixes
- V1 vs. V2, OSS vs. Cloud vs. Enterprise

7 Mogelijke examenvragen

- Wat zijn multi stage builds? Wat is het nut?
 - Het opsplitsen van een dockerfile voor optimalisatie en leesbaarheid
 - Verhoogt leesbaarheid en optimaliseert de image (zodat het sneller gaat)
- Wat is het voordeel van docker-compose?
 - Zorgt voor een leesbare file waar alle images samen staan
 - Je kan alle images in één keer opstarten
 - Goed voor security: alle containers zijn isolated
- Wat zijn cgroups?
 - Control groups == een type namespace
 - Limits, accounts for & isolates resource usage
- Wat is een image?

- Een file bestaande uit meerdere layers dat gebruikt worden om code uit te voeren in een container
 - Een file met instructie regels over hoe de container moet gebouwd worden
- Wat is een index uit wat bestaat het?
 - Index heeft een logische namespace waar in staat welke mapping, shards en replica's het heeft
 - gelijkaardig aan een database in RDBMS
- Wat is het nut van mapping en wat is het?
 - Mapping geeft structure aan de dataset in Elastic Stack
 - We geven datatypes mee en zeggen hoe de verdeling van clusters en shards is
- Wat is een measurements?
 - Een string die informatie geeft over de associated fields.
 - Data structuur in InfluxDB
- Wat is een buffer?
 - Een buffer houdt data vast voor het verwerkt wordt
 - de data kan zo asynchroon worden ingelezen
- Wat is de tick stack?
 - Telegraf: to collect and write data, metrics and events using plugins
 - InfluxDB: Time-series database
 - Chronograf: real-time visualisation of InfluxDB data
 - Kapacitor: monitoring & alerting based on views of InfluxDB data and anomalies
- Wat is een lambda functie:
 - een container die weinig resources gebruikt
 - deze kan events ontvangen en kan gaan reageren op een bepaald event
- Wat is AWS lambda:
 - Serverless compute service
 - Code zonder managing servers
 - ⇒ scripting service die automatisch je script/container uitvoert en managed
- Wat is multi factor auth (MFA) en wat is het nut:
 - Elektronische authenticatie methode (2FA)
 - Een user kan pas binnen als hij 2 fases van authenticatie kan doorstaan
 - 2 dingen: iets dat de user heeft en iets dat de user weet ! ENKEL DE USER
- Wat is het verschil tussen expose & ports:
 - expose laat enkel de uitvoerende container op de poort
 - ports laat alle connecties toe op die poort