

MLOps

Tuur Vanhoutte

September 27, 2021

Contents

1	Introduction	1
1.1	Why do I need deployment automation	1
1.2	What will we do	1
1.3	The AI pipeline	1
1.4	Teleport	1
1.4.1	Teleport Server Access	2
2	FastAPI	2
3	Docker	2
4	Kubernetes	2
4.1	What is Kubernetes?	2
4.2	Architecture	2
4.2.1	Control Plane / Kubernetes Master	3
4.2.2	Node	3
4.3	Objects	4
4.3.1	Building blocks	4
4.3.2	Pods	4
4.3.3	Services	5
4.3.4	Deployments and ReplicaSets	5
4.3.5	Daemonsets	5
4.3.6	StatefulSets	5
4.3.7	PersistentVolumes and PersistentVolumeClaims	5
4.3.8	ConfigMaps and Secrets	6
4.4	Load balancer	6
4.5	Autoscaler	7
4.6	YAML file	7
4.7	Helm	8
5	Rancher	8
5.1	What we learned	10

1 Introduction

- Deployment of deep learning models in an API
- Deploying services to Kubernetes with Docker
- Version Controlling with Git (to blame those who ruined the code)
- Automation pipelines to prepare data, train models and deploy applications
 - Automated delivery (ready to deploy)
 - Automated deployment (CI / CD)

1.1 Why do I need deployment automation

- Save time
- Increased accuracy
- Better documentation
- Autoscaling

1.2 What will we do

- Object classification with a CNN
- Deployed with FastAPI
- Scalable with Kubernetes
- Automated with CI/CD pipelines
 - Yes, also automated AI Training!

1.3 The AI pipeline

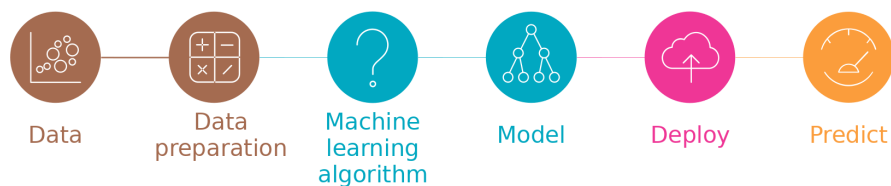


Figure 1: The complete AI pipeline

1.4 Teleport

How to connect to your dedicated virtual machine? **Teleport** allows engineers and security professionals to unify access for SSH servers, Kubernetes clusters, web applications, and databases across all environments.

1.4.1 Teleport Server Access

- One centralized login and user-management system
- Randomly generated passwords for your user account
- We manage all virtual machines your use can access

To **install and use Teleport**, see Leho.

2 FastAPI

TODO

3 Docker

TODO

4 Kubernetes

4.1 What is Kubernetes?

Definition 4.1 *Kubernetes is a **portable, extensible, open-source** platform for **managing containerized workloads and services**, that facilitates both **declarative configuration** and **automation**. It has a large, rapidly growing ecosystem. Kubernetes services, support, and tools are widely available.*

- Service discovery and load balancing
- Storage orchestration
- Automated rollouts and rollbacks
- Self-healing
- Auto-scaling




Figure 2: Kubernetes workflow

4.2 Architecture

A Kubernetes cluster consists of a set of worker machines, called nodes, that run containerized applications. Worker nodes host the Pods that are the components of the application workload. The control plane manages the worker nodes and the Pods in the cluster. In production environments,

the control plane usually runs across multiple computers and a cluster usually runs multiple nodes, providing fault-tolerance and high availability.

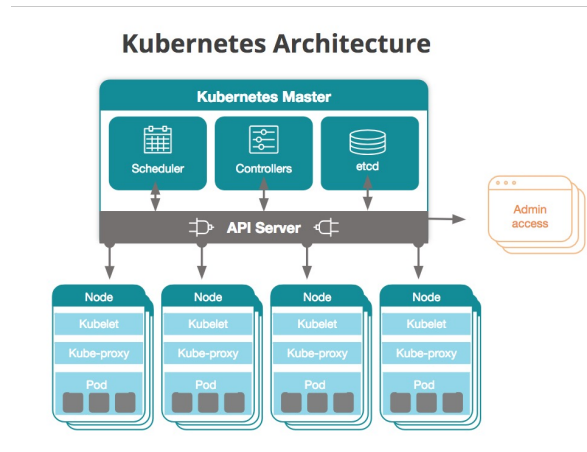


Figure 3: Architecture

4.2.1 Control Plane / Kubernetes Master

- Scheduler
 - Makes sure that the number of desired pods is always running
 - Keeps track of capacity and resources of nodes
 - Assigns work to nodes based on availability
- Controller
 - responsible for registering nodes and monitoring health
- etcd
 - A persistent and distributed key-value data store
 - Stores the state and configuration data for the entire cluster
- API-server
 - main access point to the control plane and master node

4.2.2 Node

A node is a virtual machine. It has at least two tasks:

- Kubelet
 - Manages the state of a node
 - * Starting, stopping and maintaining application containers with instructions from the control plane
 - Collects performance and health information of node, pods and containers

- * Shares those with the Control Plane
- Kube-proxy
 - Network proxy
 - Manages the Virtual IP addresses of pods and services
 - Also works as a load balancer for services running on a node
- Pod

4.3 Objects

Kubernetes objects are persistent entities in the Kubernetes system. Kubernetes uses these entities to represent the state of your cluster. Specifically, they can describe:

- What containerized applications are running (and on which nodes)
- The resources available to those applications
- The policies around how those applications behave, such as restart policies, upgrades, and fault-tolerance

A Kubernetes object is a "record of intent"—once you create the object, the Kubernetes system will constantly work to ensure that object exists. By creating an object, you're effectively telling the Kubernetes system what you want your cluster's workload to look like; this is your cluster's desired state.

4.3.1 Building blocks

- **Pods**
- **Services** and EndPoints
- **Deployments**
- ReplicaSets
- DaemonSets
- StatefulSets
- **Ingress**
- PersistentVolumes and PersistentVolumeClaims
- ConfigMaps and Secrets

4.3.2 Pods

Definition 4.2 *In Kubernetes, the **Pod** serves as a kind of basic, functional unit. A pod is a set of containers, along with their shared volumes.*

As Kubernetes' scheduler creates and deletes application pods unexpectedly, you should not rely on a particular pod. However, you do need to be able to access your application in a predictable manner. And to do that, Kubernetes provides the simplest form of load balancing traffic, namely a Service.

4.3.3 Services

Definition 4.3 *A Kubernetes Service is an abstraction which defines a logical set of Pods and a policy by which to access them - sometimes called a micro-service. The set of Pods targeted by a Service is (usually) determined by a Label Selector.*

Services have IP addresses (used internally by Kubernetes) which are relatively stable. When a program element needs to make use of the functions abstracted by the service, it makes a request to the service, rather than an individual pod. The service then acts as a dispatcher, assigning a pod to handle the request. Thus, a client never connects to a container, but to a Pod, through a Service.

4.3.4 Deployments and ReplicaSets

Definition 4.4 *A deployment is a YAML object that defines the pods and the number of container instances, called replicas, for each pod. You define the number of replicas you want to have running in the cluster via a ReplicaSet.*

Definition 4.5 *A ReplicaSet is part of the deployment object. So, for example, if a node running a pod dies, the replica set will ensure that another pod is scheduled on another available node.*

4.3.5 Daemonsets

Definition 4.6 *A DaemonSet ensures that all (or some) Nodes run a copy of a Pod. As nodes are added to the cluster, Pods are added to them. As nodes are removed from the cluster, those Pods are garbage collected. Deleting a DaemonSet will clean up the Pods it created.*

Some typical uses of a DaemonSet are:

- running a cluster storage daemon, such as glusterd, ceph, on each node.
- running a logs collection daemon on every node, such as fluentd or filebeat.

4.3.6 StatefulSets

Definition 4.7 *Like a Deployment, a StatefulSet manages Pods that are based on an identical container spec. Unlike a Deployment, a StatefulSet maintains a sticky identity for each of their Pods. These pods are created from the same spec, but are not interchangeable: each has a persistent identifier that it maintains across any rescheduling.*

If you want to use storage volumes to provide persistence for your workload, you can use a StatefulSet as part of the solution. Although individual Pods in a StatefulSet are susceptible to failure, the persistent Pod identifiers make it easier to match existing volumes to the new Pods that replace any that have failed.

4.3.7 PersistentVolumes and PersistentVolumeClaims

PersistentVolumes

- Define a storage volume in the cluster
- Independent lifecycle
- Otherwise: ephemeral data inside the pods

4.3.8 ConfigMaps and Secrets

- Decouple configurations from hard-coded environment variables
- Separate Kubernetes object to share with Pods
- Difference: Secrets hold confidential information
- Example:
 - Different configurations for Development, Staging and Production environments, quickly interchangeable
 - Storing passwords and using it in your applications

4.4 Load balancer

L4 Load Balancing (TCP)

- Only available on Kubernetes Cloud Providers (GCP, AWS, Azure ...)

L7 Load Balancing

- Can redirect traffic to specific workloads based on request
- Ingress

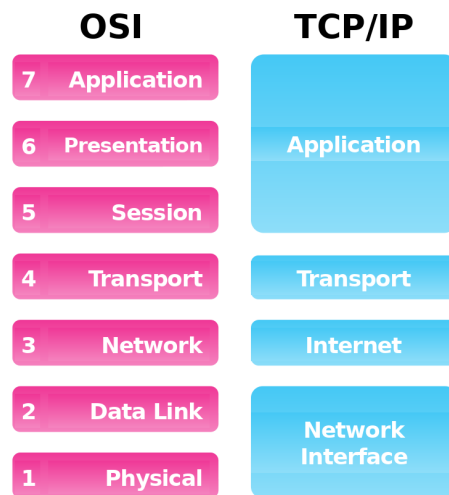


Figure 4: Layers in OSI & TCP/IP

- Kube-proxy
 - Load distribution
 - Default: Randomly route to IP addresses based on iptables
 - Previously: Round-robin
- Ingress
 - Routes traffic based on request rules configured
 - Deploy a Controller (like nginx)

- Deploy Resources
- Tip: Use Rancher to deploy Ingress Controllers and Resources (it has a UI for that)

4.5 Autoscaler

Kubernetes uses the horizontal pod autoscaler (HPA) to monitor the resource demand and automatically scale the number of replicas. When changes in replica count are required, the number of replicas is increased or decreased accordingly.

You define the minimum and maximum number of replicas that can run. You also define the metric to monitor and base any scaling decisions on, such as CPU usage.




Figure 5: Autoscaler

4.6 YAML file

```

1  apiVersion: apps/v1
2  kind: Deployment
3  metadata:
4    name: nginx-deployment
5  spec:
6    selector:
7      matchLabels:
8        app: nginx
9    replicas: 2 # tells deployment to run 2 pods
10   template:
11     metadata:
12       labels:
13         app: nginx
14     spec:
15       containers:
16       - name: nginx
17         image: nginx:1.14.2
18         ports:
19         - containerPort: 80

```

- apiVersion
 - Which version of the Kubernetes API you're using to create this object.
- kind
 - What kind of object you want to create

- metadata
 - Data that helps uniquely identify the object, including the name string, UID and optional namespace
- spec
 - What state you desire for the object

4.7 Helm

Definition 4.8 *Helm is an application package management registry for Kubernetes, maintained by the CNCF. Helm "charts" are pre-configured software application resources you can download and deploy and in your Kubernetes environment.*

According to a 2018 CNCF survey, 68% of respondents said Helm was the preferred package management tool for Kubernetes applications. Helm charts can help DevOps teams come up to speed more quickly with managing applications in Kubernetes; it allows them leverage existing charts that they can share, version, and deploy into their dev and production environments.

- Package manager
- We can create custom Helm charts

```

1 spec:
2   containers:
3   - args:
4     - uwsgi
5     - --ini
6     - app.ini
7   env:
8     {{- range $key, $value := .Values.env.webapp }}
9     - name: {{ $key }}
10      value: {{ $value | quote }}
11    {{- end }}
12   image:{{ .Values.repository }}
13   /basicwebapp:{{ .Values.app.version.webapp }}
14   name: {{ $name }}
15   resources: {}
16   restartPolicy: Always

```

- Organised in template.yaml files
- Placeholders inside yaml files
- Values.yaml to fill placeholders
- Installing and upgrading versions of a release through CLI
- Loops, conditionals, variables, filters, values and more ...

5 Rancher

Definition 5.1 *Rancher is a complete software stack for teams adopting containers.*

It addresses the operational and security challenges of managing multiple Kubernetes clusters, while providing DevOps teams with integrated tools for running containerized workloads.

- Simple, intuitive UI to get started with Kubernetes
- Multi-cloud and hybrid-cloud Kubernetes solutions
- Fast way to set up on-premises Kubernetes clusters
- Includes CI/CD pipelines for DevOps operations

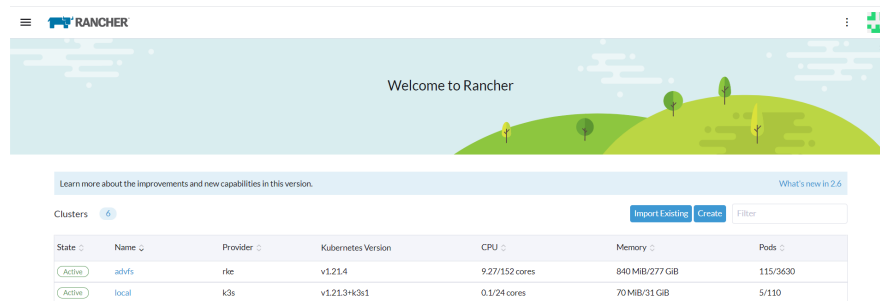


Figure 6: Rancher welcome page

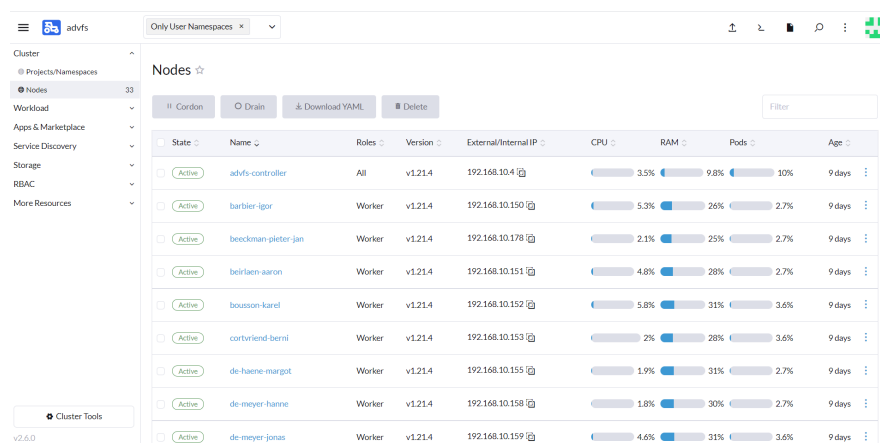


Figure 7: List of nodes

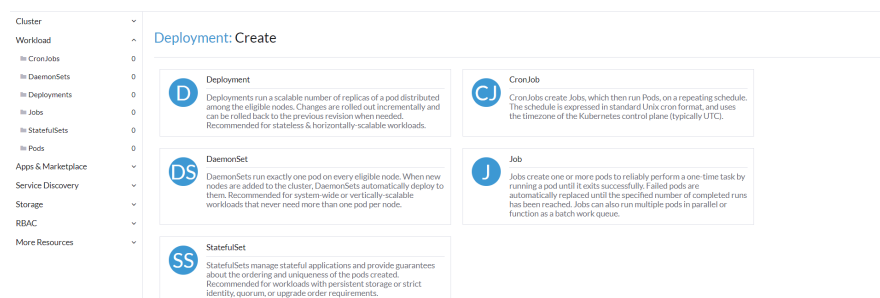


Figure 8: Deployment

Figure 9: Creating a Deployment

Figure 10: Creating a Deployment

5.1 What we learned

- Why Kubernetes is useful for microservices
- The Kubernetes Building Blocks
- How autoscaling in Kubernetes works
- Different kinds of Load Balancing options
- How to write a YAML file for a Kubernetes object
- What the advantages of Helm charts are
- What Rancher does for us

Setting up Rancher and Kubernetes: see labs